

Per-token, per-layer gradient L2 norm — contradiction (N=20, no CoT, full attention), 1207 tokens, loss 0.92

